Complete eBOSS P(k) analysis: (1) Load DR14 observational data, (2) Compare CDM, wCDM, and CDM+massive neutrinos models with data in a plot, (3) Run MCMC for , m, N_eff with Planck priors on the massive neutrino model, (4) Show posterior corner plot and best-fit estimates.

HEP Multi-agent System
Argonne Computational Sciences Division

February 8, 2026

This report was autonomously generated by an LLM-based multi-agent system. All content, analysis, and conclusions should be independently verified.

1 Query

Run everything via MCP tools in mcp-ke. Absolutely do not use your own code (no writing new python codes, existing tools should be used). Strictly no fake/synthetic/placeholder data. If something doesn't work, do not try to find a workaround or shortcuts that need writing realistic estimations.

(1) First, load observational data from eBOSS (see: /data/a/cpac/nramachandra/Projects/AmSC/mcp-ke/data/DR14_pm3d_19kbins.txt). (2) Then compare the P(k) with wCDM, CDM + Massive Neutrinos and CDM (use any set of parameters you need). Plot them all. (3) Finally, run a full posterior analysis using MCMC? Do this for 3 parameters (sigma8, m, N_species) of the CDM + massive neutrinos model. Use priors from Planck. Go with a small run. (4) Show me the final posterior distribution plot from GetDist and the best-fit estimates.

2 Answer

The agentic system successfully completed the first two stages of the requested eBOSS power spectrum analysis: (1) loading the DR14 observational data (19 k-bins) and (2) computing theoretical $P(k)$ predictions for three cosmological models (ΛCDM, wCDM, and ΛCDM + massive neutrinos) using CLASS with Planck 2018 best-fit parameters at $z_{\rm eff} = 1.48$. However, the pipeline stalled at the comparison plot step due to a serialization bug in the `plot_power_spectra` MCP tool, which prevented all downstream tasks—including the MCMC posterior analysis, corner plot, and best-fit estimates—from executing. No MCMC results, posterior distributions, or best-fit parameter constraints were obtained.

3 Execution Log

Step 1: Load Observational Data (Completed)

The compute worker invoked `load_observational_data` on the file `DR14_pm3d_19kbins.txt`, which contains 19 rows and 3 columns (no header). The tool parsed the columns as k [h/Mpc], $P(k)$ [(Mpc/h)^3], and $\sigma_{P(k)}$ [(Mpc/h)^3]. Three individual `.npy` arrays (`k_obs.npy`, `Pk_obs.npy`, `Pk_obs_err.npy`) and a metadata JSON file were saved via `save_array` and `save_dict`. Verification was performed by reloading with `load_array`; all 19 bins matched the source file. No errors occurred.

Step 2: Create Theory k-Grid (Completed)

The compute worker called `create_theory_k_grid` to produce a 300-point logarithmic grid from 10^{-4} to 10 h/Mpc. The grid was saved as `theory_k_grid.npy` and verified by round-trip loading. No errors occurred.

Step 3: Compute All Three Model Spectra (Completed)

The compute worker first attempted `compute_all_models` with manually assembled parameter dictionaries. This failed because the parameter keys did not match CLASS conventions (`w0` instead of the fluid approximation flag; `sum_mnu_eV` and `N_species` are not CLASS parameters). The worker then retrieved canonical parameter sets via three helper tools:

- `get_lcdm_params` — returned the standard Planck 2018 six-parameter set.
- `get_wcdm_params` — returned the same base set plus `_w0_approx = -0.9`.
- `get_nu_mass_params` — returned the base set with 3 degenerate non-cold dark matter species ($\Sigma m_\nu = 0.15\,\mathrm{eV}$, $N_{\mathrm{ncdm}} = 3$, $N_{\mathrm{ur}} = 0.044$).

Each parameter dictionary was modified to set `z_pk = 1.48` before being passed individually to `compute_power_spectrum`. All three CLASS runs succeeded. The resulting arrays (each 300 points) were saved as `Pk_lcdm.npy`, `Pk_wcdm.npy`, and `Pk_nu_mass.npy`, along with `model_params.json`. The worker verified the outputs by computing model ratios.

Step 4: Comparison Plot (Failed)

The visualization worker loaded all arrays and attempted `plot_power_spectra` multiple times with varying input formats:

1. Passing 300-element k-theory list: failed with *"x and y must have same first dimension, but have shapes (1,) and (300,)"*.
2. Passing 19-element arrays (matching the observational grid): failed with the identical shape mismatch error.
3. Passing a dictionary reference: failed with *"unhashable type: dict"*.
4. Invoking the orchestrating `power_spectrum_agent`: failed due to missing environment variables (`LLM_API_KEY`, `LLM_URL`, `LLM_MODEL`).

The worker logged the issue twice via `log_issue` and exhausted the 25-iteration retry limit without producing a plot. The root cause is a deserialization bug in the tool: the JSON array parameter for `k_theory` is wrapped as `np.array([list])`, producing shape $(1, N)$ instead of $(N,)$.

Steps 5–9: MCMC, Corner Plot, Posterior Analysis, Best-Fit $P(k)$, Citations (All Skipped)

Because Step 4 blocked the pipeline, no MCMC sampling (`run_mcmc_cosmology`), corner plot (`create_mcmc_corner_plot`), posterior summary (`analyze_mcmc_samples`), best-fit spectrum (`compute_best_fit_power_spectrum`), or citation retrieval was executed.

4 Results

Observational Data

The eBOSS DR14 quasar power spectrum contains **19 k-bins** spanning $k =$ **0.208–2.521 h/Mpc** at effective redshift $z_{\text{eff}} =$ **1.48**. The measured $P(k)$ ranges from ~ 3.4 **to** ~ 804 **(Mpc/h)^3**, with uncertainties from ~ 1.9 **to** ~ 263 **(Mpc/h)^3**.

Theoretical Power Spectra

All models were evaluated on a 300-point log-spaced grid at $z = 1.48$ using CLASS. Table **??** summarizes the key properties.

Table 1: Theoretical $P(k)$ summary for three cosmological models at $z_{\text{eff}} = 1.48$.

Model	Peak $P(k)$ [(Mpc/h)^3]	k_{peak} [h/Mpc]	Ratio to ΛCDM ($k\sim 0.1$)	Ratio to ΛCDM ($k\sim 1.0$)	Key parameter
ΛCDM (Planck 2018)	136,591	0.0059	1.000	1.000	—
wCDM ($w_0 = -0.9$)	126,041	0.0059	0.923	0.923	$w_0 = -0.9$
ΛCDM + ν ($\Sigma m_\nu = 0.15\,\text{eV}$)	124,179	0.0057	0.857	0.855	$N_{\text{ncdm}} = 3$

The wCDM model shows a uniform **∼7.7%** suppression relative to ΛCDM across all scales, consistent with modified late-time structure growth from a less negative dark energy equation of state. The massive neutrino model exhibits a larger, mildly scale-dependent **∼14.3%** suppression, with slightly stronger damping at high k due to neutrino free-streaming.

MCMC Posterior Analysis

No results were produced. The MCMC run targeting σ_8, Σm_ν, and N_{eff} with Planck priors was never executed due to the upstream plotting failure.

5 Discussion

Physical Plausibility. The computed model power spectra exhibit physically expected behavior. The peak of $P(k)$ near $k \sim 0.006\,\text{h/Mpc}$ corresponds to the matter–radiation equality scale, consistent with Planck 2018 cosmology at $z = 1.48$. The wCDM suppression is scale-independent to within numerical precision, as expected for a modified background expansion history that uniformly alters the linear growth factor. The massive neutrino suppression of ∼14% for $\Sigma m_\nu = 0.15\,\text{eV}$ is consistent with the standard approximation $\Delta P/P \approx -8 f_\nu$ where $f_\nu = \Omega_\nu/\Omega_m$; for $\Sigma m_\nu = 0.15\,\text{eV}$ and $\Omega_m h^2 \approx 0.143$, $f_\nu \approx 0.011$, giving ∼9% suppression at the linear level. The slightly larger observed suppression (∼14%) at $z = 1.48$ is plausible given the combined effects of altered background evolution and the redshift dependence of neutrino free-streaming.

Limitations. The analysis is severely incomplete. Without the comparison plot, no visual validation of model–data agreement was performed. Without MCMC, no parameter constraints, posterior distributions, or best-fit estimates exist. The theoretical spectra use the matter power spectrum $P_m(k)$, whereas the eBOSS data may require a bias model, shot noise subtraction, and window function convolution for rigorous comparison—none of which were applied.

Gaps. The wCDM implementation uses a phenomenological flag (`_w0_approx = -0.9`) whose mapping to the CLASS `fluid` or `ppf` dark energy module is unclear from the tool output. The parameter helper tools return $z_{\text{pk}} = 0$ by default; the worker manually overrode this to $z = 1.48$, but the documentation does not confirm whether additional redshift-dependent settings (e.g., nonlinear corrections) were needed.

Failed Steps. The `plot_power_spectra` tool failure is the critical bottleneck. It blocked 5 of 8 downstream tasks. The bug is deterministic and reproducible: every invocation wraps the `k_theory` array into shape $(1,)$ regardless of input length.

6 Developer Notes

- `plot_power_spectra` — **Critical Bug (Array Deserialization).** The tool fails with ``` ``x and y must have same first dimension, but have shapes (1,) and (N,)'' ``` for any input array size (tested with $N = 19$ and $N = 300$). Root cause: the MCP JSON interface deserializes the `k_theory` parameter as `np.array([list_of_values])`, producing shape `(1, N)` or even `(1,)` instead of `(N,)`. **Fix:** Add `k_theory = np.asarray(k_theory).flatten()` at the entry point of the tool's implementation. Additionally, the same flattening should be applied to all array-type parameters (`pk_values` in model results, `k_obs`, `pk_obs`, `pk_obs_err`). This bug renders the tool completely unusable via the JSON/MCP interface and must be fixed before any plotting workflow can succeed.

- `plot_power_spectra` — **Secondary Bug (Dict Input).** When `model_results` is passed as a nested dictionary (model name → array), the tool raises ``` ``unhashable type: dict'' ```. This suggests the tool attempts to use the dictionary as a hashable key or passes it to a function expecting a scalar. The input parsing logic for `model_results` should be reviewed.

- `power_spectrum_agent` — **Missing Environment Variables.** The orchestrating agent requires `LLM_API_KEY`, `LLM_URL`, and `LLM_MODEL` environment variables that are not set in the MCP server configuration. This agent cannot be used as a fallback for the broken plotting tool.

- `compute_all_models` — **Parameter Key Mismatch.** When called with user-supplied parameter dictionaries containing keys like `w0`, `sum_mnu_eV`, or `N_species`, CLASS raises `CosmoSevereError` because these are not valid CLASS input parameters. The tool should either (a) document the exact CLASS parameter names required or (b) implement internal key translation from common aliases. The workaround of calling `get_*_params` helpers first is effective but undocumented.

- **Pipeline Architecture.** The sequential dependency chain meant that a single tool failure (plotting) cascaded into 5 skipped tasks, including the entire MCMC analysis. Consider restructuring the pipeline so that compute-heavy tasks (MCMC) do not depend on visualization steps.